



AI AGENTS PLUS
AI development & automation

CLIENT PROPOSAL

GO AI Platform

Autonomous Website Generation & Operations Platform

A connected operating system that turns a prospect enquiry into a personalised proposal and a hosted, on-brand website — AI for speed, a human for quality. Engineered by AI Agents Plus.

PREPARED FOR
Shane Kearley

DATE
June 2026

PREPARED BY
AI Agents Plus

PROPOSAL REF
AAP-PROP-2026-014

The Opportunity

Go AI's vision is an AI-first digital agency that takes a small business from first enquiry to a live, premium website with as little manual effort as possible — starting in Corfu and Greece, then scaling out. The ambition is clear and the tooling exists. What is missing is the operating system that connects intake, proposal, build and delivery into one pipeline, so nothing is re-entered by hand and quality never slips as volume grows.

This engagement builds exactly that. The platform captures a structured brief from each prospect, generates and emails a personalised proposal, and hands a clean brief to the Hermes generation agent, which builds the customer site from an approved template library, deploys it to Vercel, and routes it through a human approval gate before anything is published. Every step — every lead, proposal, email, payment and build — is tracked in a single operations dashboard, with the agent reporting its own actions back so the whole business is visible in one place.

It is engineered around the one thing that makes or breaks an AI website business: the output has to look genuinely premium, not generated. The Hermes agent builds only from a curated design system, and a person signs off on every site before a customer ever sees it. That is the layer that protects both quality and margin as Go AI scales toward fifty customers and beyond.

At a glance

- Prospect enquiry to hosted website through one connected pipeline
- Hermes agent builds from an approved, premium template library
- Personalised proposals generated and emailed automatically via Resend
- Human approval gate before any site is published — quality never automated away
- Operations dashboard tracking every lead, proposal, email, build and customer
- Built on Next.js, Supabase, Vercel and Stripe to scale to 50+ customers

Why this matters now

Warm leads are already in hand through the Corfu gyms, with a target of fifty customers in six months.

Building the engine before outbound lead generation begins means the system is ready the moment volume arrives — not retrofitted under pressure once the business is already moving.

Three Principles

The platform is built on three convictions that keep it premium as it scales.

01

Premium, never generated

Every site is assembled from an approved design system and curated components, not a blank prompt each time. The output has to earn the word premium — that is the entire reason Go AI exists.

02

AI for speed, humans for quality

Agents handle the repetitive work at scale — intake, proposals, drafts, builds — but a person signs off on every site before a customer sees it. Speed without a quality gate is how AI businesses lose trust.

03

One brief, one source of truth

A single structured customer record powers the proposal, the build, the dashboard and the reporting. Nothing is re-entered, nothing drifts out of sync, and the system can grow in stages without costly rewrites.

What we'll build

Eight connected modules that take a prospect from enquiry to a hosted, on-brand website.

01 Quote Intake & Structured Brief

A guided intake captures business type, requirements, package, assets and brand details, validates the inputs, and writes a clean structured brief that every other module reads from.

02 AI Proposal Generation

From the brief, the platform drafts a personalised one-page proposal — business summary, recommended package, add-ons, pricing and next steps — ready to go out.

03 Automated Email Delivery

Proposals and notifications are sent automatically through Resend, and every send is logged back to the dashboard as a tracked event against that customer.

04 Hermes Website-Generation Agent

On a confirmed brief, the Hermes agent builds the customer site from approved templates, generates the copy and sections, and produces a per-customer repository and a live preview.

05 Approved Design System & Templates

A curated set of premium templates, components and design tokens the agent must build from — the layer that keeps every output on-brand and away from the generic AI look.

06 Human Review & Publish Gate

Each generated site arrives as a preview with a build summary for sign-off. Nothing is published to a customer without explicit human approval.

07 Hosting & Deployment

Approved sites deploy to Vercel with SSL and a connected domain, ready to hand over and host as part of the recurring model.

08 Operations Dashboard

One control centre tracking every lead, proposal, email, payment, build and customer status — with the Hermes agent reporting its actions live so the whole pipeline is visible at a glance.

Technical Architecture

A modern, scalable stack chosen for fast iteration and long-term growth.

APPLICATION & DASHBOARD

Next.js · React · TypeScript

A fast, reliable application and operations dashboard on one modern framework, typed end to end for maintainability.

DATA LAYER

Supabase

Database, authentication and the central customer record — the single source of truth every module reads from and writes to.

GENERATION AGENT

Hermes Agent + Claude Code / Design

Orchestrates the build from approved GitHub template repositories, generating each customer site under a controlled, repeatable process.

DESIGN SYSTEM

Tailwind + Tokenised Components

A curated component library and design tokens the agent assembles from, keeping every output premium and consistent.

DEPLOYMENT

Vercel

Per-customer deployment, preview URLs, SSL and domain connection for fast, reliable hosting.

EMAIL DELIVERY

Resend

Transactional proposal and notification email, with every send captured as a dashboard event for live client tracking.

PAYMENTS

Stripe

One-off setup and recurring services, with payment status written straight to the customer record.

REPORTING & EVENTS

Event-driven dashboard

Agent and system actions stream into the dashboard so every client and every step is trackable in real time.

Quality, Security & Compliance

A product that touches user data carries a responsibility we take seriously. The standards below are not afterthoughts — they are baked into the engineering process from the first commit.

Code quality

Clean, layered architecture with separation of concerns. Peer-reviewed pull requests and automated tests on every commit. Production code that resists rot and adapts to future requirements.

Security

Sensitive data encrypted at rest and in transit. No credentials in source control. Security headers, input validation, and dependency scanning baked into the CI pipeline.

Performance & reliability

Measured performance budgets enforced in CI. Production monitoring with anomaly alerts. Designed for graceful degradation under load and intermittent connectivity.

Privacy & compliance

GDPR principles applied by default: data minimisation, user-controlled export and deletion, and no third-party tracking on first-party experiences.

Accessibility

WCAG 2.1 AA targets where applicable. Dynamic type sizing, high-contrast mode support, and semantic labelling for assistive technologies.

Ownership & handover

Full source code, deployment runbooks, and architectural documentation transferred at handover. No vendor lock-in, no dependency on us to operate the system.

An Eight-Week Build

Four two-week milestones, each ending in something reviewable and signed off before the next begins.

MILESTONE 1

01

Milestone 1

Foundation & Customer Brief

Weeks 1 – 2

- Project setup, repository structure and environments.
- Database, authentication and the central customer-record design.
- The core data architecture every module builds on.
- Reviewable: admin login and customer-record schema in place.

MILESTONE 2

02

Milestone 2

Intake & AI Proposal Engine

Weeks 3 – 4

- Guided quote and intake workflow writing a clean structured brief.
- AI proposal generation from the brief — spec, packages, pricing.
- Automated proposal email via Resend, logged to the customer record.
- Reviewable: enquiry in, personalised proposal out by email.

MILESTONE 3

03

Milestone 3

Hermes Generation Agent

Weeks 5 – 6

- Approved template and component library set up as the build source.
- Hermes agent: brief to generated site, copy and a per-customer repo.
- Vercel deployment with a preview URL and the human review gate.
- Reviewable: a brief produces a previewable, on-brand site.

MILESTONE 4

04

Milestone 4

Dashboard, Reporting & Automation

Weeks 7 – 8

- Operations dashboard: leads, proposals, payments, builds, statuses.
- Hermes agent reporting its actions live into the dashboard.
- Stripe billing, automation foundation, full testing and refinements.
- Reviewable: production-ready MVP handed over.

What we'll need from you

Delivery on the agreed timeline depends on a small number of inputs from your side. Each item below is something we cannot proceed without — flagged up front to avoid surprises.

01 Vercel access

Team access to the Vercel account used as the deployment target for generated sites.

02 Email domain & Resend

A Resend account and DNS access on the sending domain for proposal and notification email.

03 Stripe account

An active Stripe account for the one-off and recurring billing integration.

04 GitHub access

A GitHub organisation for the template library and the per-customer repositories.

05 Brand & seed templates

Go AI brand assets and the starting reference designs that seed the approved template library.

06 Sign-off authority

A single decision-maker available for milestone reviews and approvals to keep the schedule on track.

What could go wrong, and how we handle it

Generated sites drifting toward a generic look

MITIGATION

The agent builds only from the approved template library, and every site passes a human review gate before publishing.

Scope expanding across the wider nine-phase vision

MITIGATION

This engagement is fixed to the platform scope below; lead generation, avatar and AI follow-up are scoped and priced as later phases.

Third-party limits or downtime (Vercel, Resend, Stripe, LLM)

MITIGATION

Retries and queued processing, graceful degradation, and clear status surfaced in the dashboard.

Outbound lead scraping being unreliable or non-compliant

MITIGATION

Deferred by design — warm leads are already in hand; when built, it runs on a licensed business-data API rather than raw scraping.

How we work together

Visibility and short feedback loops are the difference between an engagement that ships on time and one that drifts. The rhythms below are what you can expect from day one.

CADENCE

Weekly progress updates

End-of-week summary of what's shipped, what's in flight, and what's next. Live demos whenever there's something visual to walk through.

CODE ACCESS

GitHub collaborator

All code committed to a private GitHub repository from Day 1. You're added as a collaborator with full read access — visibility into every commit, no black boxes.

CHANNELS

Slack / WhatsApp / Email

Direct line via the channel that fits your workflow. Typical response within a few business hours during the working day in EAT (UTC+3).

ESCALATION

Direct to founder

No account managers, no handoffs. Single decision-maker on each side, fast turnaround on any blocker that needs an urgent call.

Platform investment

A fixed price for the full platform build, paid across four milestones. Premium customer websites built on the platform are billed separately at \$750 each.

MILESTONE 01 Foundation & Customer Weeks 1 – 2 \$1,000	MILESTONE 02 Intake & AI Proposal Weeks 3 – 4 \$1,250	MILESTONE 03 Hermes Generation Weeks 5 – 6 \$1,500	MILESTONE 04 Dashboard, Reporting & Weeks 7 – 8 \$1,250
TOTAL INVESTMENT			\$5,000

Each milestone on sign-off, via Fiverr. Premium sites \$750 each.

What happens after launch

The work doesn't end at handover. Three options below cover everything from a no-cost first-month warranty to a long-term retainer — pick the level that fits the way you want to operate.

WARRANTY

First 30 days,
included

- ✓ Bug fixes on delivered scope
- ✓ Adjustments to agreed behaviour
- ✓ No additional charge

PREMIUM SITES

\$750 per website

- ✓ Each customer site built on the platform
- ✓ Designed and assembled to the premium standard
- ✓ Billed per build

ON-DEMAND

Quoted per request

- ✓ New features and integrations
- ✓ Later phases — avatar, follow-up, lead gen
- ✓ Scoped and priced individually, no lock-in

Why AI Agents Plus

AI Agents Plus builds premium, AI-powered products end to end — design, front-end and automation — to the quality bar Go AI is built on.

1,500+

hours of advanced AI development experience powering modern, intelligent applications

30+

client projects delivered across Fiverr, Upwork, and direct engagements

100%

code ownership transferred to every client — no lock-in, no vendor dependency

How We Move Forward

- 1 Review**
Review this proposal and share any adjustments so the scope fits Go AI exactly.
- 2 Sign-off**
On acceptance, the engagement is set up through Fiverr against the four milestones.
- 3 Kick-off**
A kick-off call to finalise access, brand assets and the seed templates for the library.
- 4 Build & deliver**
Milestone-by-milestone delivery, each one reviewed and signed off before the next begins.



Let's build it.

Shane — the engine is ready to build when you are.

GET IN TOUCH

WEBSITE

ai-agentsplus.com

EMAIL

hello@ai-agentsplus.com

LOCATION

Nairobi, Kenya